

9

Markov Chains: Introduction

We now start looking at the material in Chapter 4 of the text. As we go through Chapter 4 we'll be more rigorous with some of the theory that is presented either in an intuitive fashion or simply without proof in the text.

Our focus is on a class of *discrete-time* stochastic processes. Recall that a discrete-time stochastic process is a sequence of random variables $\mathbf{X} = \{X_n : n \in I\}$, where I is a discrete index set. Unless otherwise stated, we'll take the index set I to be the set of nonnegative integers $I = \{0, 1, 2, \dots\}$, so

$$\mathbf{X} = \{X_n : n = 0, 1, 2, \dots\}.$$

We'll denote the state space of \mathbf{X} by S (recall that the state space is the set of all possible values of any of the X_i 's). The state space is also assumed to be discrete (but otherwise general), and we let $|S|$ denote the number of elements in S (called the *cardinality* of S). So $|S|$ could be ∞ or some finite positive integer.

We'll start by looking at some of the basic structure of Markov chains.

Markov Chains:

A discrete-time stochastic process X is said to be a *Markov Chain* if it has the *Markov Property*:

Markov Property (version 1):

For any $s, i_0, \dots, i_{n-1} \in S$ and any $n \geq 1$,

$$P(X_n = s | X_0 = i_0, \dots, X_{n-1} = i_{n-1}) = P(X_n = s | X_{n-1} = i_{n-1}).$$

In words, we say that the distribution of X_n given the entire past of the process only depends on the immediate past. Note that we are not saying that, for example X_{10} and X_1 are independent. They are not. However, given X_9 , for example, X_{10} is conditionally independent of X_1 . Graphically, we may imagine being on a particle jumping around in the state space as time goes on to form a (random) sample path. The Markov property is that the distribution of where I go to next depends only on where I am now, not on where I've been. This property is a reasonable assumption for many (though certainly not all) real-world processes. Here are a couple of examples.

Example: Suppose we check our inventory once a week and replenish if necessary according to some schedule which depends on the level of the stock when we check it. Let X_n be the level of the stock at the beginning of week n . The Markov property for the sequence $\{X_n\}$ should be reasonable if it is reasonable that the distribution of the (random) stock level at the beginning of the following week depends on the stock level now (and the replenishing rule, which we assume depends only on the current stock level), but not on the stock level from previous weeks. This assumption would be true if it were true that demands for the stock for the coming week do not depend on the past stock levels. □

Example: In the list model example, suppose we let X_n denote the list ordering after the n th request for an item. If we assume (as we did in that example) that all requests are independent of one another, then the list ordering after the next request should not depend on previous list orderings if I know the current list ordering. Thus it is natural here to assume that $\{X_n\}$ has the Markov property. \square

Note that, as with the notion of independence, in applied modeling the Markov property is not something we usually try to prove mathematically. It usually comes into the model as an *assumption*, and its validity is verified either empirically by some statistical analysis or by an underlying a priori knowledge about the system being modeled. The bottom line for an applied model is how well it ends up predicting an aspect of the real system, and in applied modeling the “validity” of the Markov assumption is best judged by this criterion.

A useful alternative formulation of the Markov property is:

Markov Property (version 2):

For any $s, i_0, i_1, \dots, i_{n-1} \in S$ and any $n \geq 1$ and $m \geq 0$

$$P(X_{n+m} = s | X_0 = i_0, \dots, X_{n-1} = i_{n-1}) = P(X_{n+m} = s | X_{n-1} = i_{n-1})$$

In words, this says that the distribution of the process at any time point in the future given the *most recent* past is independent of the *earlier* past. We should prove that the versions of the Markov property are equivalent, because version 2 appears on the surface to be more general. We do this by showing that each implies the other. It's clear that version 2 implies version 1 just by setting $m = 0$. We can use conditioning and an induction argument to prove that version 1 implies version 2, as follows.

Proof that version 1 implies version 2: Version 2 is certainly true for $m = 0$ (it is exactly version 1 in this case). The induction hypothesis is to assume that version 2 true holds for some arbitrary fixed m and the induction argument is to show that this implies it must also hold for $m + 1$. If we condition on X_{n+m} then

$$\begin{aligned} &P(X_{n+m+1} = s | X_0 = i_0, \dots, X_{n-1} = i_{n-1}) \\ &= \sum_{s' \in S} P(X_{n+m+1} = s | X_{n+m} = s', X_0 = i_0, \dots, X_{n-1} = i_{n-1}) \\ &\quad \times P(X_{n+m} = s' | X_0 = i_0, \dots, X_{n-1} = i_{n-1}). \end{aligned}$$

For each term in the sum, for the first probability we can invoke version 1 of the Markov property and for the second probability we can invoke the induction hypothesis, to get

$$\begin{aligned} &P(X_{n+m+1} = s | X_0 = i_0, \dots, X_{n-1} = i_{n-1}) \\ &= \sum_{s' \in S} P(X_{n+m+1} = s | X_{n+m} = s', X_{n-1} = i_{n-1}) \\ &\quad \times P(X_{n+m} = s' | X_{n-1} = i_{n-1}) \end{aligned}$$

Note that in the sum, in the first probability we left the variable X_{n-1} in the conditioning. We can do that because it doesn't affect the distribution of X_{n+m+1} conditioned on X_{n+m} . The reason we leave X_{n-1} in the conditioning is so we can use the basic property that

$$P(A \cap B | C) = P(A | B \cap C) P(B | C)$$

for any events A , B and C (you should prove this for yourself if you don't quite believe it). With $A = \{X_{n+m+1} = s\}$, $B = \{X_{n+m} = s'\}$ and $C = \{X_{n-1} = i_{n-1}\}$, we have

$$\begin{aligned} &P(X_{n+m+1} = s | X_0 = i_0, \dots, X_{n-1} = i_{n-1}) \\ &= \sum_{s' \in S} P(X_{n+m+1} = s, X_{n+m} = s' | X_{n-1} = i_{n-1}) \\ &= P(X_{n+m+1} = s | X_{n-1} = i_{n-1}). \end{aligned}$$

So version 2 holds for $m + 1$ and by induction it holds for all m . \square

Time Homogeneity: There is one further assumption we will make about the process \mathbf{X} , in addition to the Markov property. Note that we have said that the Markov property says that the distribution of X_n given X_{n-1} is independent of X_0, \dots, X_{n-2} . However, that doesn't rule out the possibility that this (conditional) distribution could depend on the time index n . In general, it could be, for example, that the distribution of where I go next given that I'm in state s at time 1 is *different* from the distribution of where I go next given that I'm in (the same) state s at time 10. Processes where this could happen are called *time-inhomogeneous*. We will assume that the process is *time-homogeneous*. Homogeneous means "same" and time-homogeneous means "the same over time". That is, every time I'm in state s , the distribution of where I go next is the same. You should check to see if you think that time homogeneity is a reasonable assumption in the last two examples.

Mathematically, what the time homogeneity property says is that

$$P(X_{n+m} = j | X_m = i) = P(X_{n+m+k} = j | X_{m+k} = i)$$

for any $i, j \in S$ and any m, n, k such that the above indices are nonnegative. The above conditional probabilities are called the *n -step transition probabilities* of the chain because they are the conditional probabilities of where you will be after n time units from where you are now. Of basic importance are the *1-step transition probabilities*:

$$P(X_n = j | X_{n-1} = i) = P(X_1 = j | X_0 = i).$$

We denote these 1-step transition probabilities by p_{ij} , and these do not depend on the time n because of the time-homogeneity assumption.

The 1-Step Transition Matrix: We think of putting the 1-step transition probabilities p_{ij} into a matrix called the 1-step transition matrix, also called the transition probability matrix of the Markov chain. We'll usually denote this matrix by \mathbf{P} . The (i, j) th entry of \mathbf{P} (i th row and j th column) is p_{ij} . Note that \mathbf{P} is a square $(|S| \times |S|)$ matrix (so it could be infinitely big). Furthermore, since

$$\sum_{j \in S} p_{ij} = \sum_{j \in S} P(X_1 = j | X_0 = i) = 1,$$

each row of \mathbf{P} has entries that sum to 1. In other words, each row of \mathbf{P} is a probability distribution over S (indeed, the i th row of \mathbf{P} is the conditional distribution of X_n given that $X_{n-1} = i$). For this reason we say that \mathbf{P} is a *stochastic matrix*.

It turns out that the transition matrix \mathbf{P} gives an almost complete mathematical specification of the Markov chain. This is actually saying quite a lot. In general, we would say that a stochastic process was specified mathematically once we specify the state space and the joint distribution of any subset of random variables in the sequence making up the stochastic process. These are called the *finite-dimensional distributions* of the stochastic process. So for a Markov chain that's quite a lot of information we can determine from the transition matrix \mathbf{P} .

One thing that is relatively easy to see is that the 1-step transition probabilities determine the n -step transition probabilities, for any n . This fact is contained in what are known as the *Chapman-Kolmogorov Equations*.

The Chapman-Kolmogorov Equations:

Let $p_{ij}(n)$ denote the n -step transition probabilities:

$$p_{ij}(n) = P(X_n = j | X_0 = i)$$

and let $\mathbf{P}(n)$ denote the n -step transition probability matrix whose (i, j) th entry is $p_{ij}(n)$. Then

$$\mathbf{P}(m+n) = \mathbf{P}(m)\mathbf{P}(n)$$

which is the same thing as

$$p_{ij}(m+n) = \sum_{k \in S} p_{ik}(m)p_{kj}(n).$$

In words: the probability of going from i to j in $m+n$ steps is the sum over all k of the probability of going from i to k in m steps, then from k to j in n steps.

Proof. By conditioning on X_m we have

$$\begin{aligned} p_{ij}(m+n) &= P(X_{m+n} = j | X_0 = i) \\ &= \sum_{k \in S} P(X_{m+n} = j | X_m = k, X_0 = i) P(X_m = k | X_0 = i) \\ &= \sum_{k \in S} P(X_{m+n} = j | X_m = k) P(X_m = k | X_0 = i) \\ &= \sum_{k \in S} P(X_n = j | X_0 = k) P(X_m = k | X_0 = i) \\ &= \sum_{k \in S} p_{kj}(n) p_{ik}(m), \end{aligned}$$

as desired. □

As we mentioned, contained in the Chapman-Kolmogorov Equations is the fact that the 1-step transition probabilities determine the n -step transition probabilities. This is easy to see, since by repeated application of the Chapman-Kolmogorov Equations, we have

$$\begin{aligned} \mathbf{P}(n) &= \mathbf{P}(n-1)\mathbf{P}(1) \\ &= \mathbf{P}(n-2)\mathbf{P}(1)^2 \\ &\vdots \\ &= \mathbf{P}(1)\mathbf{P}(1)^{n-1} \\ &= \mathbf{P}(1)^n. \end{aligned}$$

But $\mathbf{P}(1)$ is just \mathbf{P} , the 1-step transition probability matrix. So we see that $\mathbf{P}(n)$ can be determined by raising the transition matrix \mathbf{P} to the n th power. One of the most important quantities that we are interested in from a Markov chain model are the *limiting probabilities*

$$\lim_{n \rightarrow \infty} p_{ij}(n) = \lim_{n \rightarrow \infty} P(X_n = j | X_0 = i).$$

As an aside, we mention here that based on the Chapman-Kolmogorov Equations, we see that one way to consider the limiting probabilities is to examine the limit of \mathbf{P}^n . Thus, for finite dimensional \mathbf{P} , we could approach the problem of finding the limiting probabilities purely as a linear algebra problem. In fact, for any (finite) stochastic matrix \mathbf{P} , it can be shown that it has exactly one eigenvalue equal to 1 and all the other eigenvalues are strictly less than one in magnitude. What does this mean? It means that \mathbf{P}^n converges to a matrix where each row is the eigenvector of \mathbf{P} corresponding to the eigenvalue 1.

However, we won't pursue this argument further here, because it doesn't work when \mathbf{P} is infinite-dimensional. We'll look at a more probabilistic argument that works for any \mathbf{P} later. But first we need to look more closely at the structure of Markov chains.

10

Classification of States

Let us consider a fairly general version of a Markov chain called the random walk.

Example: Random Walk. Consider a Markov chain \mathcal{X} with transition probabilities

$$p_{i,i+1} = p_i,$$

$$p_{i,i} = r_i,$$

$$p_{i,i-1} = q_i \text{ and}$$

$$p_{ij} = 0 \text{ for all } j \neq i-1, i, i+1,$$

for any integer i , where $p_i + r_i + q_i = 1$. A Markov chain like this, which can only stay where it is or move either one step up or one step down from its current state, is called a *random walk*. Random walks have applications in several areas, including modeling of stock prices and modeling of occupancy in service systems. The random walk is also a great process to study because it is relatively simple to think about yet it can display all sorts of interesting behaviour that is illustrative of the kinds of things that can happen in Markov Chains in general.

To illustrate some of this behaviour, we'll look now at two special cases of the random walk. The first is the following

Example: *Random Walk with Absorbing Barriers.* Let N be some fixed positive integer. In the random walk, let

$$\begin{aligned} p_{00} &= p_{NN} = 1, \\ p_{i,i+1} &= p \text{ for } i = 1, \dots, N-1, \\ p_{i,i-1} &= q = 1 - p \text{ for } i = 1, \dots, N-1, \text{ and} \\ p_{ij} &= 0 \text{ otherwise.} \end{aligned}$$

We can take the state space to be $S = \{0, 1, \dots, N\}$. The states 0 and N are called *absorbing barriers*, because if the process ever enters in either of these states, the process will stay there forever after. \square

This example illustrates the four possible relationships that can exist between any pair of states i and j in any Markov chain.

1. State j is *accessible* from state i but state i is not accessible from state j . This means that, starting in state i , there is a positive probability that we will eventually go to state j , but starting in state j the probability that we will ever go to state i is 0. In the random walk with absorbing barriers, state N is accessible from state 1, but state 1 is not accessible from state N .
2. State j is not accessible from state i but i is accessible from j .
3. Neither states i nor j are accessible from the other. In the random walk with absorbing barriers, states 0 and N are like this.
4. States i and j are accessible from each other. In the random walk with absorbing barriers, i and j are accessible from each other for $i, j = 1, \dots, N-1$. In this case we say that states i and j *communicate*.

Accessibility and Communication: The more mathematically precise way to say that state j is accessible from state i is that the n -step transition probabilities $p_{ij}(n)$ are strictly positive for at least one n . Conversely, if $p_{ij}(n) = 0$ for all n , then state j is not accessible from state i .

The relationship of communication is our first step in classifying the states of an arbitrary Markov chain, because this relationship is what mathematicians call an *equivalence relationship*. What this means is that

- State i communicates with itself for every state i . This is true by convention, as

$$P(X_n = i | X_0 = i) = 1 > 0 \quad \text{for any state } i.$$

- If state i communicates with state j then state j communicates with state i . This is true directly from the definition of communication.
- If state i communicates with state j and state j communicates with state k , then state i communicates with state k . This is easily shown. Please see the text for a short proof.

The fact that communication is an equivalence relationship allows us to divide up the state space of any Markov chain in a natural way. We say that the *equivalence class* of state i is the set of states that state i communicates with. By the first property of an equivalence relationship, we see that the *equivalence class* of any state i is not empty: it contains at least the state i . Moreover, by the third property of an equivalence relationship, we see that *all states in the equivalence class of state i must communicate with one another*.

More than that, if j is any state in the equivalence class of state i , then states i and j have the same equivalence class. You cannot, for example, have a state k that is in the equivalence class of j but not in the equivalence class of i because that means k communicates with j , and j communicates with i , but k does not communicate with i , and this contradicts the third property of an equivalence relationship. So that means we can talk about the equivalence class of states that communicate with one another, as opposed to the equivalence class of any particular state i . Now the state space of any Markov chain could have just one equivalence class (if all states communicate with one another) or it could have more than one equivalence class (you may check that in the random walk with absorbing barriers example there are three equivalence classes). But two different equivalence classes must be *disjoint* because if, say, state i belonged to both equivalence class 1 and to equivalence class 2, then again by the third property of an equivalence relationship, every state in class 1 must communicate with every state in class 2 (through communication with state i), and so every state in class 2 belongs to class 1 (by definition), and vice versa. In other words, class 1 and class 2 cannot be different unless they are disjoint.

Thus we see that the equivalence relationship of communication divides up the state space of a Markov chain into disjoint sets of equivalence classes. You may wish to go over Examples 4.11 and 4.12 in the text at this point. In the case when all the states of a Markov chain communicate with one another, we say that the state space is *irreducible*. We also say in this case that the Markov chain is irreducible. More generally, we say that the equivalence relation of communication partitions the state space of a Markov chain into disjoint, irreducible, equivalence classes.

Classification of States

In addition to classifying the relationships between pairs of states, we also can classify each individual states into one of three mutually disjoint categories. Let us first look at another special case of the random walk before defining these categories.

Example: *Simple Random Walk*. In the random walk model, if we let

$$\begin{aligned} p_i &= p, \\ r_i &= 0, \quad \text{and} \\ q_i &= q = 1 - p \end{aligned}$$

for all i , then we say that the random walk is a *simple random walk*. In this case the state space is the set of all integers,

$$S = \{\dots, -1, 0, 1, \dots\}.$$

That is, every time we make a jump we move up with probability p and move down with probability $1 - p$, regardless of where we are in the state space when we jump. Suppose the process starts off in state 0. If $p > 1/2$, your intuition might tell you that the process will, with probability 1, eventually venture off to $+\infty$, never to return to the 0 state. In this case we would say that state 0 is *transient*, because there is a positive probability that if we start in state 0 then we will never return to it.

Suppose now that $p = 1/2$. Now would you say that state 0 is transient? If it's not, that is, if we start out in state 0 and it will eventually return to state 0 with probability 1, then we say that state 0 is *recurrent*.

Transient and Recurrent States: In any Markov chain, define

$$\begin{aligned} f_i &= P(\text{Eventually return to state } i | X_0 = i) \\ &= P(X_n = i \text{ for some } n \geq 1 | X_0 = i). \end{aligned}$$

If $f_i = 1$, then we say that state i is *recurrent*. Otherwise, if $f_i < 1$, then we say that state i is *transient*. Note that every state in a Markov chain must be either transient or recurrent.

There is a pretty important applied reason why we care whether a state is transient or recurrent. Whether a state is transient or recurrent determines the kinds of questions we ask about the process. I mentioned once that the limiting probabilities $\lim_{n \rightarrow \infty} p_{ij}(n)$ are very important because these will tell us about the “long-run” or “steady-state” behaviour of the process, and so give us a way to predict how the system our process is modeling will tend to behave on average. However, if state j is transient then this limiting probability is the wrong thing to ask about because, as we’ll see, this limit will be 0.

For example, in the random walk with absorbing barriers, it is clear that every state is transient except for states 0 and N (why?). What we should be asking about is not whether we’ll be in state i , say, where $1 \leq i \leq N - 1$, at some time point far into the future, but other things, like, given that we start in state i , where $1 \leq i \leq N - 1$, a) will we eventually end up in state 0 or in state N ? or b) what’s the expected time before we end up in state 0 or state N ?

As another example, in the symmetric random walk with $p < 1/2$, suppose we start in state B , where B is some large integer. If B is a transient state, then the question to ask is not about the likelihood of being in state B in the long run, but, for example, how long before we can expect to be in state 0.

Our main result for today is about how to check whether a given state is transient or recurrent. We will follow the argument in the text which precedes Proposition 4.1 in Section 4.3. It is based on the mean number of visits to a state. In particular, if state i is recurrent then, starting in state i , state i should be visited infinitely many times. This is because, if state i is recurrent then, starting in state i , we will return to state i with probability 1. But once we are back in state i it is as if we were back at time 0 again because of the Markov property and time-homogeneity, so we are certain to come back to state i for a second time with probability 1. This process never changes. In fact, a formal induction argument in which statement n is the proposition that state i will be visited at least n times with probability 1 given that the process starts in state i , will show that statement n is true for all n . In other words, state i will be visited infinitely many times with probability 1, and so the mean number of visits to state i is infinite.

On the other hand, if state i is transient, then, starting in state i , it will return to state i with probability $f_i < 1$ and will never return with probability $1 - f_i$, and this probability of never returning to state i will be faced every time the process returns to state i . In other words, the sequence of returns to state i is equivalent to a sequence of independent Bernoulli trials, in which a “success” corresponds to never returning to state i , and the probability of success is $1 - f_i$. So the number of returns to state i is a version of the Geometric random variable that is interpreted as the number of failures until the first success in a sequence of independent Bernoulli trials. It is an easy exercise to see that when the probability of a success is $1 - f_i$, the mean number of failures until the first success is $f_i/(1 - f_i)$. This is the mean number of returns to state i if state i is transient. All we care about it at this point is that it is finite.

On the other hand, the mean number of visits to state i can be computed in terms of the n -step transition probabilities. If we define I_n to be the indicator of the event $\{X_n = i\}$ and let N_i denote the number of visits to state i , then

$$\sum_{n=0}^{\infty} I_n = \text{the number of visits to state } i = N_i.$$

Then

$$\begin{aligned} E[N_i | X_0 = i] &= E\left[\sum_{n=0}^{\infty} I_n \mid X_0 = i\right] \\ &= \sum_{n=0}^{\infty} E[I_n | X_0 = i] \\ &= \sum_{n=0}^{\infty} P(X_n = i | X_0 = i) \\ &= \sum_{n=0}^{\infty} p_{ii}(n). \end{aligned}$$

By the argument on the preceding page, state i is recurrent if and only if the mean number of visits to state i , starting in state i , is infinite. Thus the following (Proposition 4.1 in the text) has been shown:

Proposition 4.1 in text. State i is recurrent if and only if

$$\sum_{n=0}^{\infty} p_{ii}(n) = \infty.$$

This also implies that state i is transient if and only if

$$\sum_{n=0}^{\infty} p_{ii}(n) < \infty.$$

A very useful consequence of Proposition 4.1 is that we can show that recurrence is a *class property*; that is, if state i is recurrent, then so is every state in its equivalence class.

Corollary to Proposition 4.1: If states i and j communicate and state i is recurrent, then state j is also recurrent.

Proof: Since i and j communicate there exists some integers n and m such that $p_{ij}(n) > 0$ and $p_{ji}(m) > 0$. By the Chapman-Kolmogorov equations (see last lecture)

$$\mathbf{P}(m + r + n) = \mathbf{P}(m)\mathbf{P}(r)\mathbf{P}(n)$$

for any positive integer r . If we expand out the (j, j) th entry we get

$$\begin{aligned} p_{jj}(m + r + n) &= \sum_{k \in S} p_{jk}(m) \sum_{l \in S} p_{kl}(r) p_{lj}(n) \\ &= \sum_{k \in S} \sum_{l \in S} p_{jk}(m) p_{kl}(r) p_{lj}(n) \\ &\geq p_{ji}(m) p_{ii}(r) p_{ij}(n) \end{aligned}$$

by taking just the term where $k = l = i$, and where the inequality follows because every term in the sum is nonnegative. If we sum over r we get

$$\sum_{r=1}^{\infty} p_{jj}(m + r + n) \geq p_{ji}(m) p_{ij}(n) \sum_{r=1}^{\infty} p_{ii}(r).$$

But the right hand side equals infinity because $p_{ji}(m) > 0$ and $p_{ij}(n) > 0$ by choice of m and n , and state i is recurrent. Therefore, the left hand side equals infinity, and so state j is recurrent by Proposition 4.1. \square

Note that since every state must be either recurrent or transient, transience is also a class property. That is, if states i and j communicate and state i is transient, then state j cannot be recurrent because that would imply state i is recurrent. If state j is not recurrent it must be transient.

Example: In the simple random walk it is easy to see that every state communicates with every other state. Thus, there is only one class and this Markov chain is irreducible. We can classify every state as either transient or recurrent by just checking whether state 0, for example, is transient or recurrent. Especially for $p = 1/2$, it's still not clear whether state 0 is transient or recurrent. We'll leave this question till next lecture, where we will see that we need to refine our categories of transient and recurrent a little bit. \square

A very similar argument to the one in Proposition 4.1 will show that if state i is a transient state and if j is a state such that state i is accessible from state j , then the mean number of visits to state i given that the process starts in state j is finite and can be expressed as

$$E[N_i | X_0 = j] = \sum_{n=0}^{\infty} p_{ji}(n).$$

For the above sum to be convergent (that is, finite) it is necessary that $p_{ji}(n) \rightarrow 0$ as $n \rightarrow \infty$. On the other hand, if state i is not accessible from state j , then $p_{ji}(n) = 0$ for all n . Thus we have our first limit theorem, which we state as another corollary:

Corollary: If i is a transient state in a Markov chain, then

$$p_{ji}(n) \rightarrow 0 \quad \text{as } n \rightarrow \infty$$

for all $j \in S$. \square

Since we've just had a number of new definitions and results, it's prudent at this time to recap what we have done. To summarize, we categorized the relationship between any pair of states by defining the (asymmetric) notion of accessibility and the (symmetric) notion of communication. We saw that communication is an equivalence relationship and so divides up the state space of any Markov chain into equivalence classes. Next, we categorized the individual states by defining the concepts of transient and recurrent. These are fundamental definitions in the language of Markov chain theory, and should be memorized. A recurrent state is one that the chain will return to with probability 1 and a transient state is one that the chain may never return to with positive probability. Next, we proved the useful Proposition 4.1 which says that state i is recurrent if and only if $\sum_{n=0}^{\infty} p_{ii}(n)$ is equal to infinity. Otherwise, if the sum is finite, then state i is transient. Finally, we looked at two consequences of this result. The first is that recurrence and transience are class properties, and the second is that $\lim_{n \rightarrow \infty} p_{ji}(n) = 0$ if i is any transient state. We also considered two special cases of the random walk, the random walk with absorbing barriers and the simple random walk. The random walk with absorbing barriers has three equivalence classes $\{0\}$, $\{N\}$ and $\{1, \dots, N-1\}$. States 0 and N are recurrent states, while states $1, \dots, N-1$ are transient.

The simple random walk has just one equivalence class, so all states are either all transient or all recurrent. However, we have not yet determined which. Intuitively, if $p < 1/2$ or $p > 1/2$, one would guess that all states are transient. If $p = 1/2$, it's not as intuitive. What do you think? In the next lecture we will consider this problem, and also introduce the use of generating functions.

Generating Functions and the Random Walk

Today we will detour a bit from the text to look more closely at the simple random walk discussed in the last lecture. In the process I wish to introduce the notion and use of *Generating Functions*, which are similar to *Moment Generating Functions*, which you may have encountered in a previous course.

For a sequence of numbers a_0, a_1, a_2, \dots , we define the generating function of this sequence to be

$$G(s) = \sum_{n=0}^{\infty} s^n a_n.$$

This brings the sequence into another domain (where s is the argument) that is **analogous** to the spectral frequency domain that may be familiar to engineers. If the sequence $\{a_n\}$ is the probability mass function of a random variable X on the nonnegative integers (i.e. $P(X = n) = a_n$), then we call the generating function the *probability generating function* of X , and we can write it as

$$G(s) = E[s^X].$$

Generating functions have many useful properties. Here we will briefly go over two properties that will be used when we look at the random walk. The first key property is how the generating function decomposes when applied to a *convolution*. If $\{a_0, a_1, \dots\}$ and $\{b_0, b_1, \dots\}$ are two sequences then we define their convolution by the sequence $\{c_0, c_1, \dots\}$, where

$$c_n = a_0 b_n + a_1 b_{n-1} + \dots + a_n b_0 = \sum_{i=0}^n a_i b_{n-i}.$$

The generating function of the convolution is

$$\begin{aligned} G_c(s) &= \sum_{n=0}^{\infty} c_n s^n = \sum_{n=0}^{\infty} \sum_{i=0}^n a_i b_{n-i} s^n \\ &= \sum_{i=0}^{\infty} \sum_{n=i}^{\infty} a_i b_{n-i} s^n = \sum_{i=0}^{\infty} a_i s^i \sum_{n=i}^{\infty} b_{n-i} s^{n-i} \\ &= G_a(s) G_b(s), \end{aligned}$$

where $G_a(s)$ is the generating function of $\{a_n\}$ and $G_b(s)$ is the generating function of $\{b_n\}$.

In the case when $\{a_n\}$ and $\{b_n\}$ are probability distributions and X and Y are two independent random variables on the nonnegative integers with $P(X = n) = a_n$ and $P(Y = n) = b_n$, we see that the convolution $\{c_n\}$ is just the distribution of $X + Y$ (i.e. $P(X + Y = n) = c_n$). In this case the decomposition above can be written as

$$G_{X+Y}(s) = E[s^{X+Y}] = E[s^X s^Y] = E[s^X] E[s^Y] = G_X(s) G_Y(s),$$

where the factoring of the expectation follows from the independence of X and Y . In words, the generating function of the sum of independent random variables is the product of the individual generating functions of the random variables.

The second important property that we need for now is that if $G_X(s)$ is the generating function of a random variable X , then

$$G'_X(1) = E[X]$$

This is straightforward to see because

$$\begin{aligned} G'(s) &= \frac{d}{ds}(a_0 + a_1s + a_2s^2 + \dots) \\ &= a_1 + 2a_2s + 3a_3s^2 + \dots \end{aligned}$$

so that

$$G'(1) = a_1 + 2a_2 + 3a_3 + \dots = E[X],$$

assuming $a_n = P(X = n)$.

In addition, one can also easily verify that if $G(s)$ is the generating function of an arbitrary sequence $\{a_n\}$, then the n th derivative of $G(s)$ evaluated at $s = 0$ is equal to $n!a_n$. That is,

$$G^{(n)}(0) = n!a_n$$

and so

$$a_n = \frac{1}{n!}G^{(n)}(0).$$

In other words, the generating function of a sequence determines the sequence. As a special case, the probability generating function of a random variable X determines the distribution of X . This can sometimes be used in conjunction with the property on the previous page to determine the distribution of a sum of independent random variables. For example, we can use this method to show that the sum of two independent Poisson random variables again has a Poisson distribution. Now let's consider the simple random walk.

Firstly, we discuss an important property of the simple random walk process that is not shared by most Markov chains. In addition to being time homogeneous, the simple random walk has a property called *spatial homogeneity*. In words, this means that if I take a portion of a sample path (imagine a graph of the sample path with the state on the vertical axis and time on the horizontal axis) and then displace it vertically by any integer amount, that displaced sample path, conditioned on my starting value, has the same probability as the original sample path, conditioned on its starting value. Another way to say this is analogous to one way that time homogeneity was explained. That is, a Markov chain is time homogeneous because, for any state i , every time we go into state i , the probabilities of where we go next depend only on the state i and not on the time that we went into state i . For a time homogeneous process that is also spatially homogeneous, not only do the probabilities of where we go next not depend on the time we entered state i , but also do not depend on the state. For the simple random walk, no matter where we are and no matter what the time is, we always move up with probability p and down with probability $q = 1 - p$. Mathematically, we say that a (discrete-time) process is *spatially homogeneous* if for any times $n, m \geq 0$ and any displacement k ,

$$P(X_{n+m} = b | X_n = a) = P(X_{n+m} = b + k | X_n = a + k).$$

For example, in the simple random walk the probability that we are in state 5 at time 10 given that we are in state 0 at time 0 is the same as the probability that we are in state 10 at time 10 given that we are in state 5 at time 0. Together with time homogeneity, we can assert, as another example, that the probability that we ever reach state 1 given we start in state 0 at time 0 is the same as the probability that we ever reach state 2 given that we start in state 1 at time 1.

Our goal is to determine whether state 0 in a simple random walk is transient or recurrent. Since all states communicate in a simple random walk, determining whether state 0 is transient or recurrent tells us whether all states are transient or recurrent. Let us suppose that the random walk starts in state 0 at time 0. Define

$$T_r = \text{time that the walk first reaches state } r, \text{ for } r \geq 1$$

$$T_0 = \text{time that the walk first returns to state 0.}$$

Also define

$$f_r(n) = P(T_r = n | X_0 = 0)$$

for $r \geq 0$ and $n \geq 0$ (noting that $f_r(0) = 0$ for all r), and let

$$G_r(s) = \sum_{n=0}^{\infty} s^n f_r(n)$$

be the generating function of the sequence $\{f_r(n)\}$, for $r \geq 0$. Note that $G_0(1) = \sum_{n=0}^{\infty} f_0(n)$ is the probability that the walk ever returns to state 0, so this will tell us directly whether state 0 is transient or recurrent. State 0 is transient if $G_0(1) < 1$ and state 0 is recurrent if $G_0(1) = 1$. We will proceed now by first considering $G_r(s)$, for $r > 1$, then considering $G_1(s)$. We will consider $G_0(s)$ later.

To approach the evaluation of $G_r(s)$, for $r > 1$, we consider the probability $f_r(n) = P(T_r = n | X_0 = 0)$ and condition on T_1 , the time the walk first reaches state 1, to obtain via the Markov property that

$$\begin{aligned} f_r(n) &= \sum_{k=0}^{\infty} P(T_r = n | T_1 = k) f_1(k) \\ &= \sum_{k=0}^n P(T_r = n | T_1 = k) f_1(k), \end{aligned}$$

where we can truncate the sum at n because $P(T_r = n|T_1 = k) = 0$ for $k > n$ (this should be clear from the definitions of T_r and T_1). Now we may apply the time *and* spatial homogeneity of the random walk to consider the conditional probability $P(T_r = n|T_1 = k)$. By temporal and spatial homogeneity, this is the same as the probability that the first time we reach state $r - 1$ is at time $n - k$ given that we start in state 0 at time 0. That is,

$$P(T_r = n|T_1 = k) = f_{r-1}(n - k),$$

and so

$$f_r(n) = \sum_{k=0}^n f_{r-1}(n - k)f_1(k).$$

So we see that the sequence $\{f_r(n)\}$ is the convolution of the two sequences $\{f_{r-1}(n)\}$ and $\{f_1(n)\}$. Therefore, by the first property of generating functions that we considered, we have

$$G_r(s) = G_{r-1}(s)G_1(s).$$

But by applying this decomposition to $G_{r-1}(s)$, and so on, we arrive at the conclusion that

$$G_r(s) = G_1(s)^r,$$

for $r > 1$. Now we will use this result to approach $G_1(s)$. This time we condition on X_1 , the first step of the random walk, to write, for $n > 1$,

$$f_1(n) = P(T_1 = n|X_1 = 1)p + P(T_1 = n|X_1 = -1)q.$$

Now if $n > 1$, then $P(T_1 = n|X_1 = 1) = 0$ because if $X_1 = 1$ then that implies $T_1 = 1$ also, so that $T_1 = n$ is impossible. Also, again by

the time and spatial homogeneity of the random walk, we may assert that $P(T_1 = n | X_1 = -1)$ is the same as the probability that the first time the walk reaches state 2 is at time $n - 1$ given that the walk starts in state 0 at time 0. That is

$$P(T_1 = n | X_1 = -1) = f_2(n - 1).$$

Therefore,

$$f_1(n) = qf_2(n - 1),$$

for $n > 1$. For $n = 1$, $f_1(1)$ is just the probability that the first thing the random walk does is move up to state 1, so $f_1(1) = p$. We now have enough to write out an equation for the generating function $G_1(s)$. Keeping in mind that $f_1(0) = 0$, we have

$$\begin{aligned} G_1(s) &= \sum_{n=0}^{\infty} s^n f_1(n) \\ &= sf_1(1) + \sum_{n=2}^{\infty} s^n f_1(n) \\ &= ps + \sum_{n=2}^{\infty} s^n qf_2(n - 1) \\ &= ps + qs \sum_{n=2}^{\infty} s^{n-1} f_2(n - 1) \\ &= ps + qs \sum_{n=1}^{\infty} s^n f_2(n) \\ &= ps + qsG_2(s), \end{aligned}$$

since $f_2(0) = 0$ as well.

But by our previous result we know that $G_2(s) = G_1(s)^2$, so that

$$G_1(s) = ps + qsG_1(s)^2,$$

and so we have a quadratic equation for $G_1(s)$. We may write it in the more usual form

$$qsG_1(s)^2 - G_1(s) + ps = 0.$$

Using the quadratic formula, the two roots of this equation are

$$G_1(s) = \frac{1 \pm \sqrt{1 - 4pqs^2}}{2qs}.$$

Only one of these two roots is the correct answer. A *boundary condition* for $G_1(s)$ is that $G_1(0) = f_1(0) = 0$. You can check (using L'Hospital's rule) that only the solution

$$G_1(s) = \frac{1 - \sqrt{1 - 4pqs^2}}{2qs}$$

satisfies this boundary condition.

Let us pause for a moment to consider what we have done. We have seen a fairly typical way in which generating functions are used with discrete time processes, especially Markov chains. A generating function can be defined for any sequence indexed by the nonnegative integers. This sequence is often a set of probabilities defined on the nonnegative integers, and obtaining the generating function of this sequence can tell us useful information about this set of probabilities. In fact, theoretically at least, the generating function tells us everything about this set of probabilities, since the generating function determines these probabilities. In our current work with the random walk, we have defined probabilities $f_r(n)$ over the *times* $n = 0, 1, 2, \dots$ at which

an event of interest first occurs. We may also work with generating functions when the state space of our Markov chain is the set of nonnegative integers and we define probabilities over the set of *states*. Many practical systems of interest are modeled with a state space that is the set of nonnegative integers, including service systems in which the state is the number of customers/jobs/tasks in the system. We will consider the use of generating functions again when we look at *queueing* systems in Chapter 8 of the text. Note also that when the sequence is a sequence of probabilities, a typical way that we try to determine the generating function of the sequence is to use a conditioning argument to express the probabilities in the sequence in terms of related probabilities. As we have seen before, when we do this right we set up a system of equations for the probabilities, and as we have just seen now, this can be turned into an equation for the generating function of the sequence. One of the advantages of using generating functions is when we are able to compress many (usually infinitely many) equations for the probabilities into just a single equation for the generating function, as we did for $G_1(s)$.

So $G_1(s)$ can tell us something about the probabilities $f_1(n)$, for $n \geq 1$. In particular

$$\begin{aligned} G_1(1) &= \sum_{n=1}^{\infty} P(T_1 = n | X_0 = 0) \\ &= P(\text{the walk ever reaches 1} | X_0 = 0) \end{aligned}$$

Setting $s = 1$, we see that

$$G_1(1) = \frac{1 - \sqrt{1 - 4pq}}{2q}.$$

We can simplify this by replacing p with $1 - q$ in $1 - 4pq$ to get $1 - 4pq = 1 - 4(1 - q)q = 1 - 4q + 4q^2 = (1 - 2q)^2$. Since $\sqrt{(1 - 2q)^2} = |1 - 2q|$, we have that

$$G_1(1) = \frac{1 - |1 - 2q|}{2q}.$$

If $q \leq 1/2$ then $|1 - 2q| = 1 - 2q$ and $G_1(1) = 2q/2q = 1$. On the other hand, if $q > 1/2$ then $|1 - 2q| = 2q - 1$ and $G_1(1) = (2 - 2q)/2q = 2p/2q = p/q$, which is less than 1 if $q > 1/2$. So we see that

$$G_1(1) = \begin{cases} 1 & \text{if } q \leq 1/2 \\ p/q < 1 & \text{if } q > 1/2. \end{cases}.$$

In other words, the probability that the random walk ever reaches 1 is 1 if $p \geq 1/2$ and is less than 1 if $p < 1/2$.

Next we will finish off our look at generating functions and the random walk by evaluating $G_0(s)$, which will tell us whether state 0 is transient or recurrent.

12

More on Classification

Today's lecture starts off with finishing up our analysis of the simple random walk using generating functions that was started last lecture. We wish to consider $G_0(s)$, the generating function of the sequence $f_0(n)$, for $n \geq 0$, where

$$\begin{aligned} f_0(n) &= P(T_0 = n | X_0 = 0) \\ &= P(\text{walk first returns to 0 at time } n | X_0 = 0). \end{aligned}$$

We condition on X_1 , the first step of the random walk, to obtain via the Markov property that

$$f_0(n) = P(T_0 = n | X_1 = 1)p + P(T_0 = n | X_0 = -1)q.$$

We follow similar arguments as those in the last lecture to evaluate the conditional probabilities. For $P(T_0 = n | X_0 = -1)$, we may say that by the time and spatial homogeneity of the random walk, this is the same as the probability that we first reach state 1 at time $n - 1$ given that we start in state 0 at time 0. That is,

$$P(T_0 = n | X_0 = -1) = P(T_1 = n - 1 | X_0 = 0) = f_1(n - 1).$$

By a similar reasoning, if we define T_{-1} to be the first time the walk reaches state -1 and define $f_{-1}(n) = P(T_{-1} = n | X_0 = 0)$, then

$$P(T_0 = n | X_1 = 1) = P(T_{-1} = n - 1 | X_0 = 0) = f_{-1}(n - 1),$$

so we have so far that

$$f_0(n) = f_{-1}(n - 1)p + f_1(n - 1)q.$$

Now if we had $p = q = 1/2$, so that on each step the walk was equally likely to move up or down, reasoning by symmetry tells us that $P(T_{-1} = n - 1 | X_0 = 0) = P(T_1 = n - 1 | X_0 = 0)$. For general p , we can also see by symmetry that, given that we start in state 0 at time 0, the distribution of T_{-1} is the same as the distribution of T_1 in a “reflected” random walk in which we *interchange* p and q . That is, if we let $f_1^*(n)$ denote the probability that the first time we reach state 1 is at time n given that we start in state 0 at time 0, but in a random walk with p and q interchanged, then what we are saying is that

$$f_1^*(n) = f_{-1}(n)$$

So, keeping in mind what the $*$ means, we have

$$f_0(n) = f_1^*(n - 1)p + f_1(n - 1)q.$$

Therefore (since $f_0(0) = 0$),

$$\begin{aligned} G_0(s) &= \sum_{n=1}^{\infty} s^n f_0(n) = \sum_{n=1}^{\infty} s^n f_1^*(n - 1)p + \sum_{n=1}^{\infty} s^n f_1(n - 1)q \\ &= ps \sum_{n=1}^{\infty} s^{n-1} f_1^*(n - 1) + qs \sum_{n=1}^{\infty} s^{n-1} f_1(n - 1) \\ &= psG_1^*(s) + qsG_1(s), \end{aligned}$$

where $G_1^*(s)$ is the same function as $G_1(s)$ except with p and q interchanged.

Now, recalling from the last lecture that

$$G_1(s) = \frac{1 - \sqrt{1 - 4pqs^2}}{2qs},$$

we see that

$$G_1^*(s) = \frac{1 - \sqrt{1 - 4pqs^2}}{2ps},$$

and so

$$\begin{aligned} G_0(s) &= psG_1^*(s) + qsG_1(s) \\ &= ps \frac{1 - \sqrt{1 - 4pqs^2}}{2ps} + qs \frac{1 - \sqrt{1 - 4pqs^2}}{2qs} \\ &= 1 - \sqrt{1 - 4pqs^2} \end{aligned}$$

Well, that took a bit of maneuvering, but we've ended up with quite a simple form for $G_0(s)$. Now using this, it is easy to look at

$$G_0(1) = \sum_{n=1}^{\infty} f_0(n) = P(\text{walk ever returns to } 0 | X_0 = 0).$$

Plugging in $s = 1$, we get

$$G_0(1) = 1 - \sqrt{1 - 4pq}.$$

As we did before, we can simplify this by writing $1 - 4pq = 1 - 4(1 - q)q = 1 - 4q + 4q^2 = (1 - 2q)^2$, so that

$$G_0(1) = 1 - |1 - 2q|.$$

Now we can see that if $q \leq 1/2$ then $G_0(1) = 1 - (1 - 2q) = 2q$. This equals one if $q = 1/2$ and is less than one if $q < 1/2$. On the other hand, if $q > 1/2$ then $G_0(1) = 1 - (2q - 1) = 2 - 2q = 2p$, and this is less than one when $q > 1/2$.

So we see that $G_0(1)$ is less than one if $q \neq 1/2$ and is equal to one if $q = 1/2$. In other words, state 0 is transient if $q \neq 1/2$ and is recurrent if $q = 1/2$. Furthermore, since all states in the simple random walk communicate with one another, there is only one equivalence class, so that if state 0 is transient then all states are transient and if state 0 is recurrent then all states are recurrent. This should also be intuitively clear by the spatial homogeneity of the random walk. In any case, it makes sense to say that the random walk is transient if $q \neq 1/2$ and is recurrent if $q = 1/2$.

Now we remark here that we have just seen that the “random variables”, T_r , in some cases had distributions that did not sum up to one. For example, as we have just seen

$$\sum_{n=1}^{\infty} P(T_0 = n | X_0 = 0) < 1$$

if $q \neq 1/2$. This is because, for $q \neq 1/2$, it is the case that $P(T_0 = \infty | X_0 = 0) > 0$. In words, T_0 equals infinity with positive probability. According to what you should have learned in a previous probability course, this disqualifies T_0 to be called a random variable when $q \neq 1/2$. In fact, we say in this case that T_0 is a *defective* random variable. Defective random variables can arise sometimes in the study of stochastic processes, especially when the random quantity, let's call it, denotes a time until the process first does something, like reach a certain state or set of states, because it may be that the process will never reach that set of states with some positive probability. In fact, knowing that a process may never reach a state or set of states with some positive probability is of central interest for some systems, such as population models where we want to know if the population will ever die out.

In the case where $q = 1/2$, of course, T_0 is a proper random variable. Here let us go back to the second property of generating functions that we covered last lecture, which is that for a random variable X with probability generating function $G_X(s)$, the expected value of X is given by

$$E[X] = G'_X(1).$$

When $q = 1/2$, the random variable T_0 has probability generating function

$$G_0(s) = 1 - \sqrt{1 - s^2},$$

so that taking the derivative we get

$$G'_0(s) = \frac{s}{\sqrt{1 - s^2}}.$$

Setting $s = 1$ we see that $G'_0(1) = +\infty$. That is, $E[T_0] = +\infty$ when $q = 1/2$. So even though state 0 is recurrent when $q = 1/2$, we have here an illustration of a very special kind of recurrent state. Starting in state 0, we will return to state 0 with probability one, and so state 0 is recurrent. But the expected time to return to state 0 is $+\infty$. In such a case we call the recurrent state a *null recurrent state*. This is another important classification of a state which we discuss next.

Remark: The text uses a more direct approach, using Proposition 4.1, to determine the conditions under which a simple random walk is transient or recurrent. Please read Example 4.13 to see this solution. Our approach using generating functions was chosen in part to introduce generating functions, which have uses much beyond analysing the simple random walk. We were also able to obtain more information about the random walk using generating functions.

Null Recurrent and Positive Recurrent States. Recurrent states in a Markov chain can be further categorized as *null recurrent* and *positive recurrent*, depending on the expected time to return to the state. If μ_i denotes the expected time to return to state i given that the chain starts in state i , then we say that state i is null recurrent if it is recurrent (i.e. $P(\text{eventually return to state } i | X_0 = i) = 1$) but the expected time to return is infinite:

State i is null recurrent if $\mu_i = \infty$.

Otherwise, we say that state i is positive recurrent if it is recurrent and the mean time to return is finite:

State i is positive recurrent if $\mu_i < \infty$.

A null recurrent state i is, like any recurrent state, returned to infinitely many times with probability 1. Even so, the probability of being in a null recurrent state is 0 in the limit:

$$\lim_{n \rightarrow \infty} p_{ji}(n) = 0,$$

for any starting state j . This interesting fact is something we'll prove later. A very useful fact, that we'll also prove later, is that null recurrence is a class property. We've already seen that recurrence is a class property. This is stronger. Not only do all states in an equivalence class have to be recurrent if one of them is, they all have to be null recurrent if one of them is. This also implies that they all have to be positive recurrent if one of them is.

So we see that the equivalence relationship of communication divides up any state space into disjoint equivalence classes, and each equivalence class has "similar" members, in that they are either all transient, all null recurrent, or all positive recurrent.

A finite Markov chain has at least one positive recurrent state: The text (on p.192) argues that a Markov chain with finitely many states must have at least one recurrent state. The argument is basically that not all states can be transient because if this were so then eventually we would run out of states to “never return to” if there are only finitely many states. We can also show that a Markov chain with finitely many states must have at least one *positive recurrent* state, which is slightly stronger. In particular, this implies that **any finite Markov chain that has only one equivalence class has all states being positive recurrent.**

Since we must be somewhere at time n , we must have $\sum_{j \in S} p_{ij}(n) = 1$ for any starting state i . That is, each n -step transition matrix $\mathbf{P}(n)$ is a stochastic matrix (the i th row is the distribution of where the process will be n steps later, starting in state i). This is true for every n , so we can take the limit as $n \rightarrow \infty$ and get

$$\lim_{n \rightarrow \infty} \sum_{j \in S} p_{ij}(n) = 1.$$

But if S is finite, we can take the limit inside to get

$$1 = \sum_{j \in S} \lim_{n \rightarrow \infty} p_{ij}(n).$$

But if every state were transient *or* null recurrent we would have $\lim_{n \rightarrow \infty} p_{ij} = 0$ for every i and j . Thus we would get the contradiction that $1=0$. Thus, there must be at least one positive recurrent state. Keep in mind that limits cannot in general be taken inside summations if it is an infinite sum. For example, $\lim_{n \rightarrow \infty} \sum_{j=1}^{\infty} 1/n \neq \sum_{j=1}^{\infty} \lim_{n \rightarrow \infty} 1/n$. The LHS is $+\infty$ and the RHS is 0.

Recurrent Classes are Closed. Now, in general, the state space of a Markov chain could have 1 or more transient classes, 1 or more null recurrent classes, and 1 or more positive recurrent classes. While state spaces with both transient and recurrent states are of great practical use (and we'll look at these in Section 4.6), many practical systems of interest are modeled by Markov chains for which all states are recurrent, and usually all positive recurrent. Even in this case, one can imagine that there could be 2 or more disjoint classes of recurrent states. The following Lemma says that when all states are recurrent we might as well assume that there is only one class because if we are in a recurrent class we will never leave it. In other words, we say that a recurrent class is *closed*.

Lemma: Any recurrent class is closed. That is, if state i is recurrent, and state j is not in the equivalence class containing state i , then $p_{ij} = 0$.

Proof. Suppose we start the chain in state i . If p_{ij} were positive, then with positive probability we could go to state j . But once in state j we could never go back to state i because if we could then i and j would communicate. This is impossible because j is not in the equivalence class containing i , by assumption. But never going back to state i is also impossible because state i is recurrent by assumption. Therefore, it must be that $p_{ij} = 0$. \square

Therefore, if all states are recurrent (null or positive recurrent) and there is more than one equivalence class, we may as well assume that the state space consists of just the equivalence class that the chain starts in.

Period of a State:

There is one more property of a state that we need to define, and that is the *period* of a state. In words, the period of a state i is the largest integer that evenly divides all the possible times that we could return to state i given that we start in state i . If we let d_i denote the period of state i , then mathematically we write this as

$$d_i = \text{Period of state } i = \gcd\{n : p_{ii}(n) > 0\},$$

where “gcd” stands for “greatest common divisor”. Another way to say this is that if we start in state i , then we cannot return to state i at any time that is not a multiple of d_i .

Example: For the simple random walk, if we start in state 0, then we can only return to state 0 with a positive probability at times 2, 4, 6, . . . ; that is, only at even times. The greatest divisor of this set of times is 2, so the period of state 0 is 2. \square

If a state i is positive recurrent, then we will be interested in the limiting probability $\lim_{n \rightarrow \infty} p_{ii}(n)$. But before we do so we need to know the period of state i . This is because if $d_i \geq 2$ then $p_{ii}(n) = 0$ for any n that is not a multiple of d_i . But then this limit will not exist unless it is 0 because the sequence $\{p_{ii}(n)\}$ has infinitely many 0's in it. It turns out that the *subsequence* $\{p_{ii}(d_i n)\}$ will converge to a nonzero limiting value.

If the period of state i is 1, then we are fine. The sequence $\{p_{ii}(n)\}$ will have a proper limiting value. We call a state that has period 1 *aperiodic*.

Period is a Class Property. As we might have come to expect and certainly hope for, our last result for today is to show that all states in an equivalence class have the same period.

Lemma: If states i and j communicate they have the same period.

Proof. Let d_i be the period of state i and d_j be the period of state j . Since i and j communicate, there is some m and n such that $p_{ij}(m) > 0$ and $p_{ji}(n) > 0$. Using the Chapman-Kolmogorov equations again as we did on Monday, we have that

$$p_{ii}(m + r + n) \geq p_{ij}(m)p_{jj}(r)p_{ji}(n),$$

for any $r \geq 0$, and the right hand side is strictly positive for any r such that $p_{jj}(r) > 0$ (because $p_{ij}(m)$ and $p_{ji}(n)$ are both strictly positive). Setting $r = 0$, we have $p_{jj}(0) = 1 > 0$, so $p_{ii}(m + n) > 0$. Therefore, d_i divides $m + n$. But since $p_{ii}(m + r + n) > 0$ for any r such that $p_{jj}(r) > 0$ we have that d_i divides $m + r + n$ for any such r . But since d_i divides $m + n$ it must also divide r . That is, d_i divides and r such that $p_{jj}(r) > 0$. In other words, d_i is a divisor of $\{r : p_{jj}(r) > 0\}$. Since d_j is, by definition, the *greatest* common divisor of the above set of r values, we must have $d_i \leq d_j$. Now repeat the same argument but interchange i and j , to get $d_j \geq d_i$. Therefore, $d_i = d_j$. \square

By the above lemma, we can speak of the period of a class, and if a Markov chain has just one class, then we can speak of the period of the Markov chain. In particular, we will be interested for the next little while in Markov chains that are irreducible (one class), aperiodic (period 1) and positive recurrent. It is these Markov chains for which the limiting probabilities both exist and are not zero. We call such Markov chains *ergodic*.